




# Cloud Monitoring as a Service

Built On Machine Learning



## Table of Contents

- 1 Why Machine Learning
- 2 Who Cares
- 3 Four Dimensions to Cloud Monitoring
- 4 Data Aggregation
- 5 Anomaly Detection Algorithms
- 6 Preset Best Practices
- 7 Capacity Analysis
- 8 Cost Analysis
- 9 Dynamic Resource Allocation
- 10 Conclusion

# Why Machine Learning



## 1 | Don't Waste Time

Administering and scaling your own monitoring stack

Creating your own integrations

Setting thresholds one metric at a time



## 2 | Gain Extra Insights

Aggregate data by cluster and by application

See what's abnormal in your environment

Leverage industry best practices for alerting and reporting



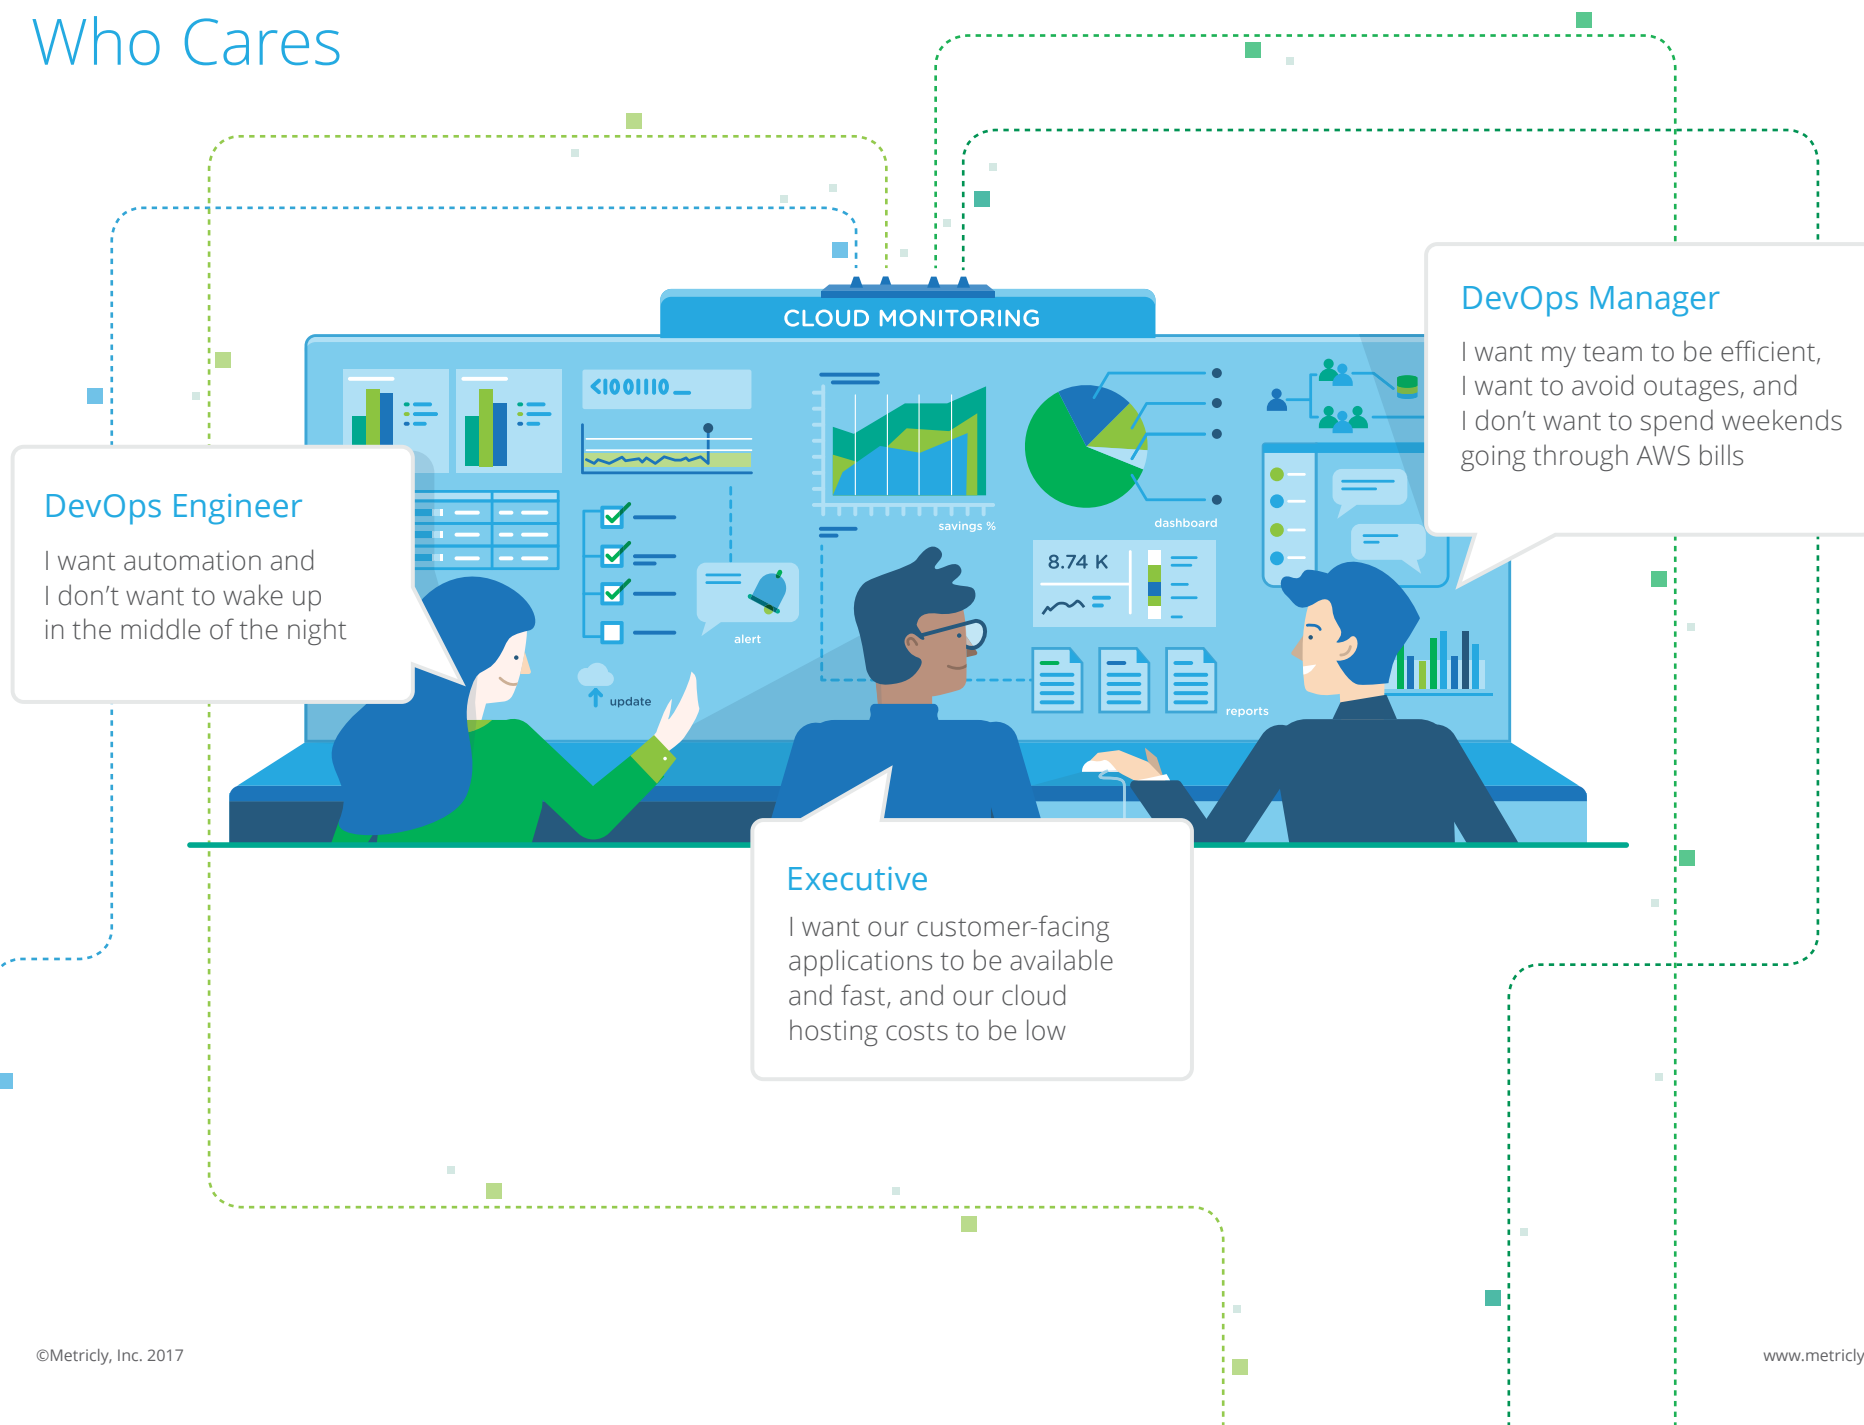
## 3 | Don't Forget Capacity & Cost

Reuse the same agent to gather capacity and performance data

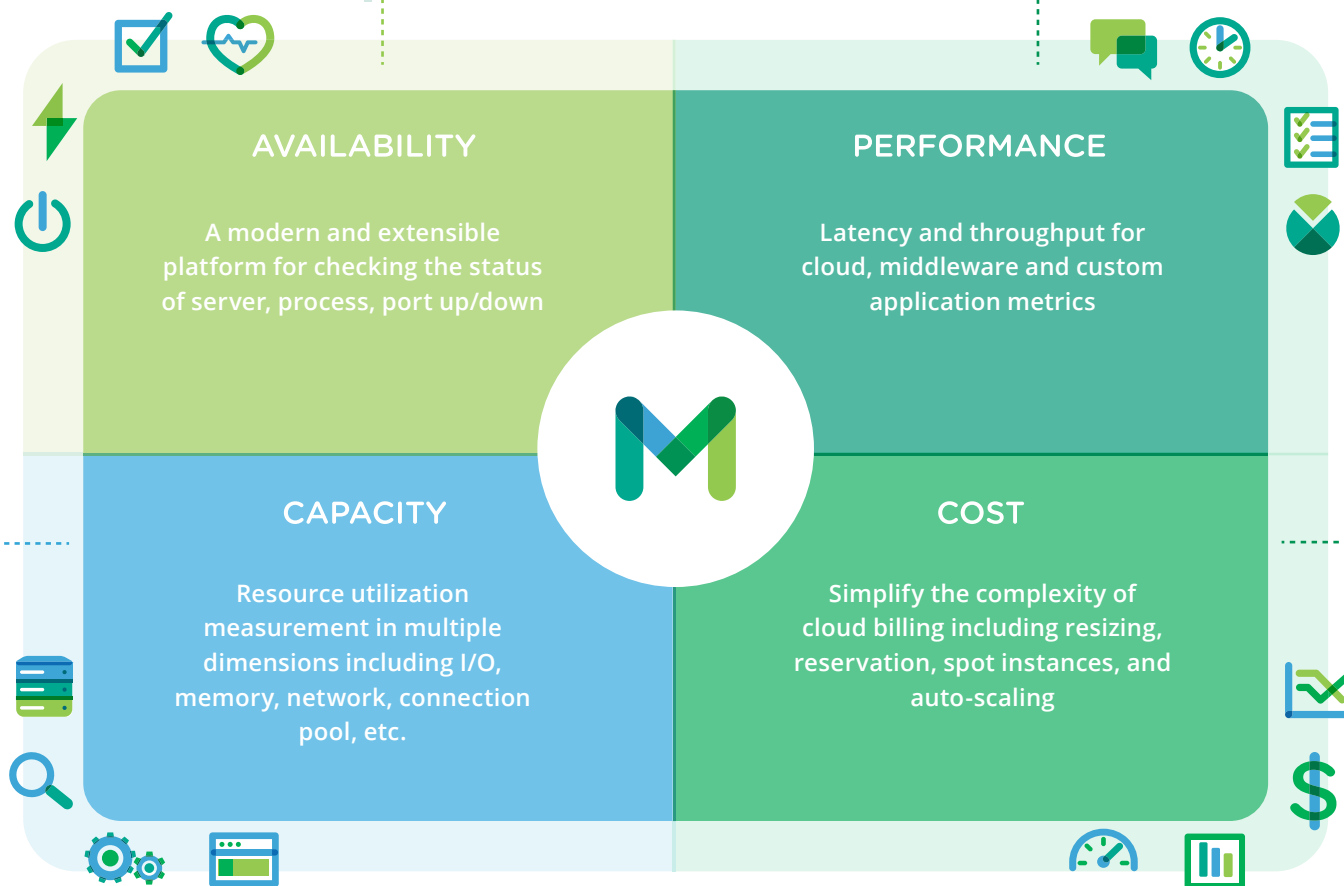
Rely on preset metrics to measure multiple dimension of capacity

AWS cost analysis gets complicated fast


# Who Cares



# Four Dimensions To Cloud Monitoring



# Data Aggregation




## A Function for Each Purpose

For an ever-increasing counter, you want to know the delta over time

**For latency, you want to know the weighed average**

For errors, you want to know the sum during an interval



## Filter and Group Metrics


User-specified metric-level aggregation is powerful

**For example, you could average CPU utilization across all EC2s with a given application tag in a particular AWS region**

Or the average latency of a method call across a cluster of Docker containers

More sophisticated expressions account for attribute meta-data such as hardware rating for IOPS for an m3.xlarge EC2 to identify unusual process queue for I/O, or the number of CPU cores that can be used to normalize the Linux load

# Anomaly Detection Algorithms





**Green Band:** Single-Variate.  
What is the value usually at this time?

**Blue Band:** Multi-Variate.  
What should it be as a function of other correlated metrics?

**Multi-variate** regression analysis discovers the regression weights and correlation coefficients to determine where a metric should be as a function of other inter-dependent metrics, for a more **accurate** estimation.

## Clusters

Anomaly detection on load-balanced clusters of nodes or containers is best accomplished by analyzing the aggregated values across the entire **cluster** vs. at the **ephemeral** node level. The aggregated values conform to the workload patterns driven by the application usage traffic.



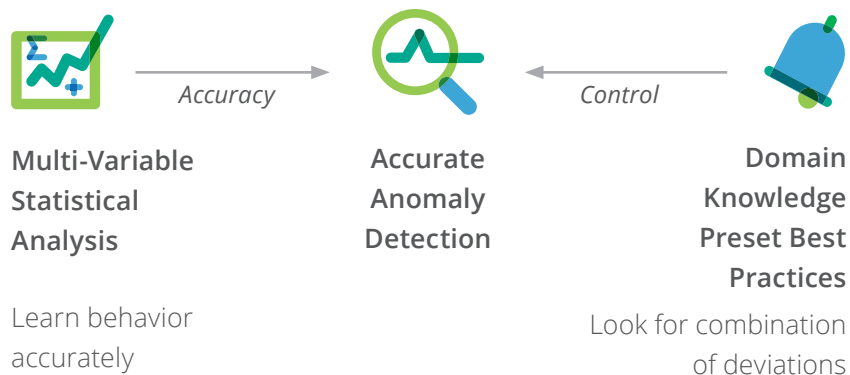
**Sudden Change:**  
Is the trend-line drastically changing up or down?

What if you don't have multiple metrics to correlate? What if your green band of normalcy is too wide due to high standard deviation? In those cases, you can detect **sudden** changes in the **trend-lines**. What's important is how steep was the rise and fall and how long it lasted.

# Preset Best Practices

## Accurate Anomaly Detection

Domain knowledge must complement statistical analysis to achieve an actionable level of accuracy in anomaly detection.



## Industry Best Practices

Preset configuration of **dashboards** and **alerts** based on industry best practices avoid the need for manual configuration by end-users upon activation of integrations with common technologies, such as the examples listed below:




## Multi-Conditional Policies


Multi-conditional alerting policies result in control, accuracy, and reasoning that are required to deliver actionable alerts. It is recommended for un-interrupted night of sleep.



1 | I/O wait > 20%



2 | Database connections drop

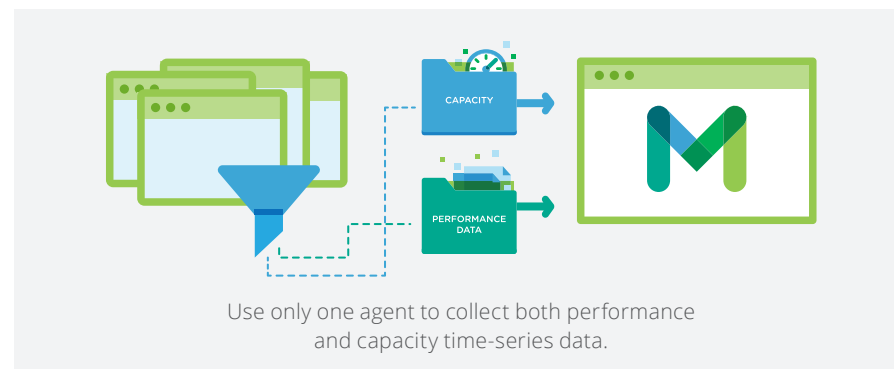


3 | ELB round-trip latency jumps




# Capacity Analysis

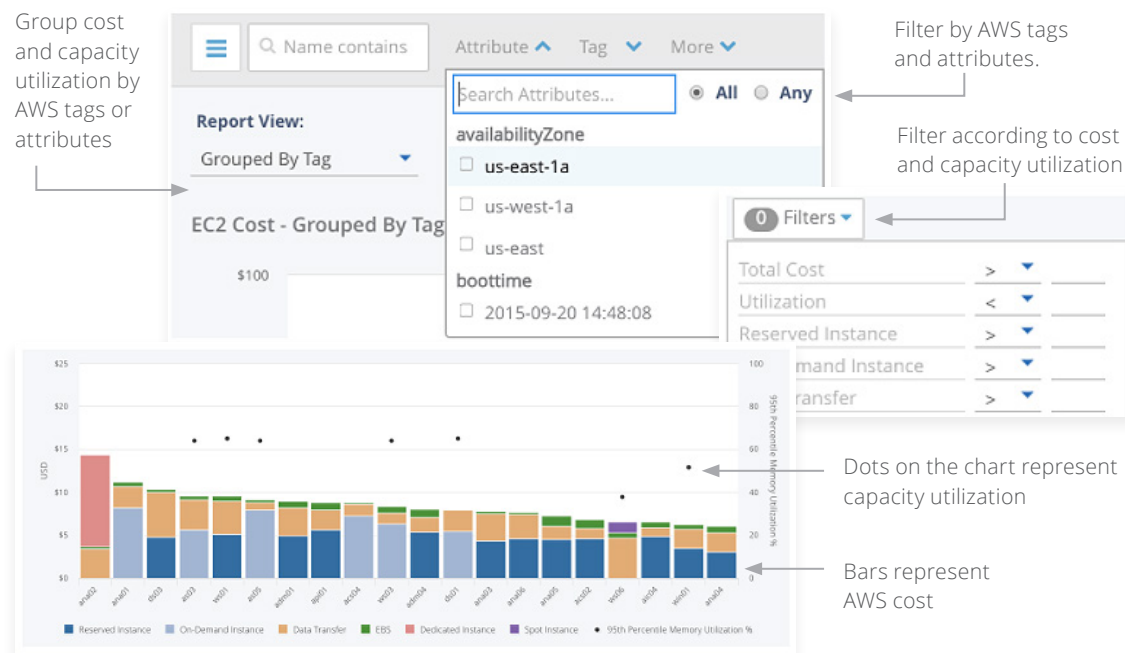
Your application platform has finite computing capacity (0 to 100%), but there are multiple types of resources. Some are **intuitive** to measure (ex. CPU). Some can be measured if you know the meta-data **attribute** (ex. 1G vs. 10G network card). Others will remain subjective, but a high water mark can help **estimate** (ex. we can process 10k concurrent users).

**CPU****MEM****NET****DISK****DB****QUEUE****LB****LATENCY****UX****BIZ**

You would miss a short-lived spike in IOPS if you take the average across a 1 or 5 minute interval. For accurate capacity planning, you must roll-up data while preserving: standard deviation, min, max, percentiles, average and median.



# Cost Analysis



## Cross-Analyze

Cross-analyze capacity utilization with AWS cost data, and use “business intelligence” oriented concepts to filter and group by meta-data. It helps you attribute cost to application’s micro-services, plan your AWS reservation by instance type, or be notified of sudden changes in utilization or cost.



**RELY ON EMAILED REPORTS TO “PUSH” VS. “PULL” SAVINGS IDEAS**

## Automate

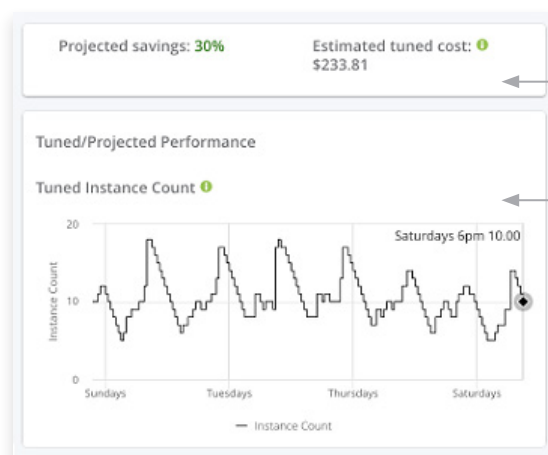
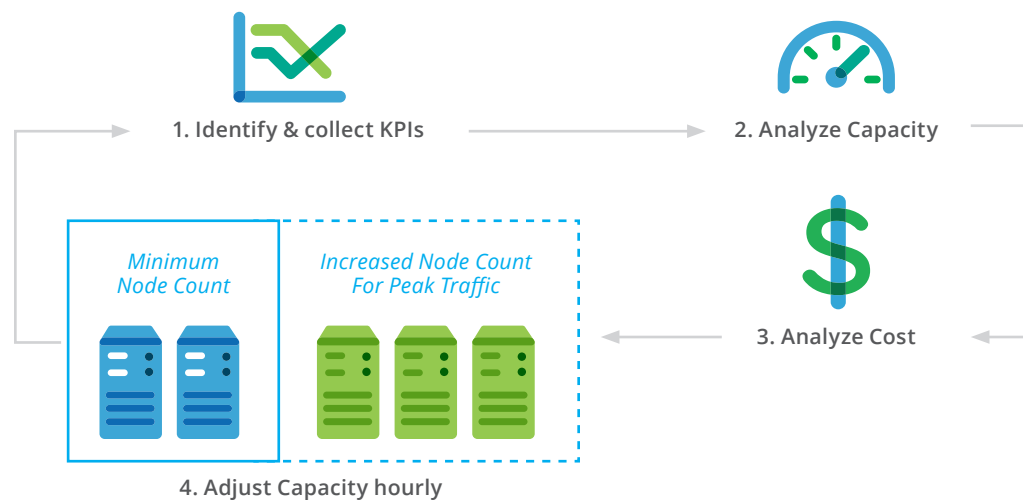
You can automate the process of searching through thousands of EC2 configurations to match the optimal EC2 configuration according to each individual instance’s workload patterns, including memory utilization and I/O.

Current Ty...	Proposed Type	Current Cost	Projected Cost	Projected Savings	Projected Savings %
m4.2xlarge	t2.medium	161.58	13.38	148.20	92%
i2.2xlarge	r3.2xlarge	200.73	79.29	121.44	61%
i2.xlarge	t2.small	117.25	4.25	113.00	96%
i2.2xlarge	r3.2xlarge	178.22	70.39	107.83	61%

# Dynamic Resource Allocation


Whether you are using Auto-Scaling Groups (ASG), Elastic Container Service (ECS), or Kubernetes, it helps to:

1. Identify Key Performance Indicators for scaling (ex. queue depth, average cluster IOPS, number of user sessions).
2. Use a tool that allows **what-if analysis** simulation based on your historic data.
3. Use the above analysis to automatically update your ASG rule-set, Cloud Formation template, or Terraform module.



Estimated cost savings compared to current configuration

Recommended hourly node count based on simulation of historic data



Simulated utilization levels across the cluster

A **what-if analysis simulator** for auto-scaling leverages your historic data to identify workload patterns and recommend optimal hourly node count according to your risk tolerance.

# Conclusion



Use Monitoring as a Service to stay focused on managing your main application stack instead of your monitoring stack...and rely on help from technical support.



Use machine learning to automate tedious manual tasks. Not only does it save time, but you will appreciate that certain types of analysis are not humanly possible.



Competition is pushing vendors to simplify and automate. Machine learning, pre-configured dashboards, and alerts make it easy to start in minutes.



Start a free trial at [www.metricly.com](http://www.metricly.com)